

Matthieu Hanania

Karis Gwet



Big Data Processing

Data Frames, Machine Learning and
NLP with Spark

Table of content

I. Movielens dataset analysis	3
A. Building an AWS cluster	3
B. Creating the movielens dataset	4
C. Apply the SQL requests	6
II. Bike rental model	10
A. Model done in class	10
B. Model improvement	13
C. The second improvement	15
D. Comparison with other machine learning model	16
III. SMS Spam Collection	18
A. Data Preparation	18
B. Training 4 models	19
C. Tuning the Random Forest model	21

I. Movielens dataset analysis

A. Building an AWS cluster

For this project, we wanted to analyze a dataset. In python, when we want to apply some treatment on a dataset it takes a lot of RAM memory. That is why, for our professional lives, we will maybe use spark instead.

Spark is a RDD system (Resilient Distributed Datasets). It means that it can execute different treatments in parallel on different clusters. It also uses the Map-Reduce algorithm. So, for Big Data projects, it is recommended to use spark.

After choosing to use Spark, we need to define clusters. We can install Spark on our computer, but we can also use a Web Service like AWS.

So, for this part, we are going to build an AWS EMR cluster.

Firstly, we created EC2 keys. They allow us to connect to our cluster. They are a file with which we can connect to the cluster using ssh protocol.

Nom

Le nom peut avoir un maximum de 255 caractères ASCII. Il ne peut pas inclure d'espaces avant ou après.

Type de paire de clés [Informations](#)

☒ RSA

☐ ED25519

Format de fichier de clé privée

☒ .pem

À utiliser avec OpenSSH

Secondly, we create a S3 bucket. It is used to save the data on the AWS server. A bucket contains an URI so we can use it to connect.

Thirdly, to execute all the treatment, we need an EMR cluster. It is composed of a DNS public address we use to connect by ssh, a master node, and two slave nodes that execute all the commands.

Récapitulatif	Détails de configuration
<p>ID : j-1RXHSRRQNQSSQ</p> <p>Date de création : 18-02-2023 22:15 (UTC+1)</p> <p>Date de fin : 18-02-2023 22:18 (UTC+1)</p> <p>Temps écoulé : 3 minutes</p> <p>Résiliation automatique : Cluster waits</p> <p>Protection de la résiliation : Désactivé</p> <p>Balises : --</p> <p>DNS public principal : ec2-15-188-53-10.eu-west-3.compute.amazonaws.com Connect to the Master Node Using SSH</p>	<p>Étiquette de version : emr-5.36.0</p> <p>Distribution Hadoop : Amazon 2.10.1</p> <p>Applications : Hive 2.3.9, Pig 0.17.0, Hue 4.10.0, JupyterHub 1.4.1, JupyterEnterpriseGateway 2.1.0, Spark 2.4.8</p> <p>URI de connexion : s3://bucket-dai-hanania/log/ 📄</p> <p>Vue cohérente EMRFS : Désactivé</p> <p>ID d'AMI personnalisée : --</p> <p>Version d'Amazon Linux : 2.0.20230119.1 En savoir plus 🔗</p>
Application user interfaces	Réseau et matériel
<p>Service d'historique : 🔗 --</p> <p>Connexions : 🔗 --</p>	<p>Zone de disponibilité : eu-west-3c</p> <p>ID de sous-réseau (subnet) : subnet-067d6a65bb107966c 🔗</p> <p>Maître : Résilié 1 m5.xlarge</p> <p>Principal : Résilié 2 m5.xlarge</p> <p>Tâche : --</p> <p>Cluster scaling: Not enabled</p>

Finally, if we want to use this cluster from home, we need to modify the VPC network to add our home IP address.

ID de règle de groupe de sécurité	Type Informations	Protocole Informations	Plage de ports Informations	Source Informations	Description - facultatif Informations
-	SSH <input type="text"/>	TCP	22	Mon IP <input type="text"/>	<input type="text"/>

After all this, we can connect to the cluster using ssh with the key and the address

```
ssh -i key_dai_Hanania.pem
hadoop@ec2-35-180-186-3.eu-west-3.compute.amazonaws.com
```

Firstly, we have to explain to the server that we want to execute pyspark commands.

```
pyspark
```

B. Creating the movielens dataset

The first dataset we have to analyze is a movielens dataset. It is composed of a “**movies**” dataset, that describes all movies (title, genres, year...) and “**ratings**” that describes each rating from user (id user, id movie, rating /5). And the concatenation of theses two datasets made the **movielens**

```
df=spark.read.csv("s3://nahle-bucket-datalake/emr/input/movielens/
movies.csv", header=True, inferSchema=True)
```

movieId	title	genres
1	Toy Story (1995)	Adventure Animation Children Comedy Fantasy
2	Jumanji (1995)	Adventure Children Fantasy
3	Grumpier Old Men (1995)	Comedy Romance

```
ratings=spark.read.csv("s3://nahle-bucket-datalake/emr/input/movie
lens/ratings.csv", header=True, inferSchema=True)
```

userId	movieId	rating	timestamp
1	307	3.5	1256677221
1	481	3.5	1256677456
1	1091	1.5	1256677471

We read these two csv files from the teacher's bucket, **header=True** means that the first line of the datasets is used to name the columns, and **inferSchema=True** attribute automatically the data types (int, string...)

After reading the two files, we can join them together. We need to import some functions that will be used later.

```
from pyspark.sql import functions as F
from pyspark.sql.functions import regexp_extract, regexp_replace
```

We want to obtain a final dataset that describes, for each movie, their mean rating, the number of rate, their title, their release year and their genres,

```
# grouping by movieId
final = ratings.groupBy("movieId").agg(F.mean('rating'),
F.count('rating'))

# joining the two datasets
final = final.join(df, ["movieId"],"inner").show()

# getting the year from the title ex : Titanic(1998)
final = final.withColumn("year", regexp_extract(final ["title"],
r"\((\d{4})\)", 1))
# removing the year from the title
```

```
final=final.withColumn("title", regexp_replace("title",
r"\\(\\d{4}\\)", ""))
```

And finally we got :

movieId	avg(rating)	count(rating)	title	genres	year
1591	2.6466656422864165	6508	Spawn	Action Adventure ...	1997
1088	3.2480141843971633	14100	Dirty Dancing	Drama Musical Rom...	1987

And we save the model in our S3 bucket. As it is a map reduce cluster, the dataset is divided into multiple subclusters.

```
final.write.csv("s3://bucket-dai-hanania/output/movielens", header
= True)
```

	Nom	Type	Dernière modification	Taille	Classe de stockage
<input type="checkbox"/>	_SUCCESS	-	15 Feb 2023 07:16:07 PM CET	0 o	Standard
<input type="checkbox"/>	part-00000-e2f6126c-a41f-476c-b618-bb0ed4a66df9-c000.csv	csv	15 Feb 2023 07:16:06 PM CET	124.5 Ko	Standard
<input type="checkbox"/>	part-00001-e2f6126c-a41f-476c-b618-bb0ed4a66df9-c000.csv	csv	15 Feb 2023 07:16:06 PM CET	126.5 Ko	Standard
<input type="checkbox"/>	part-00002-e2f6126c-a41f-476c-b618-bb0ed4a66df9-c000.csv	csv	15 Feb 2023 07:16:06 PM CET	125.3 Ko	Standard
<input type="checkbox"/>	part-00003-e2f6126c-a41f-476c-b618-bb0ed4a66df9-c000.csv	csv	15 Feb 2023 07:16:06 PM CET	123.5 Ko	Standard
<input type="checkbox"/>	part-00004-e2f6126c-a41f-476c-b618-bb0ed4a66df9-c000.csv	csv	15 Feb 2023 07:16:06 PM CET	124.0 Ko	Standard
<input type="checkbox"/>	part-00005-e2f6126c-a41f-476c-b618-bb0ed4a66df9-c000.csv	csv	15 Feb 2023 07:16:06 PM CET	128.3 Ko	Standard
<input type="checkbox"/>	part-00006-e2f6126c-a41f-476c-b618-bb0ed4a66df9-c000.csv	csv	15 Feb 2023 07:16:06 PM CET	126.5 Ko	Standard
<input type="checkbox"/>	part-00007-e2f6126c-a41f-476c-b618-bb0ed4a66df9-c000.csv	csv	15 Feb 2023 07:16:06 PM CET	125.9 Ko	Standard
<input type="checkbox"/>	part-00008-e2f6126c-a41f-476c-b618-bb0ed4a66df9-c000.csv	csv	15 Feb 2023 07:16:07 PM CET	127.0 Ko	Standard

C. Apply the SQL requests

After creating the requested dataset, we can read it again, create an SQL view and do some requests.

```
from pyspark.sql.functions import split
# read CSV
final = spark.read.option("header", "true").option("inferSchema",
"true").csv("s3://bucket-dai-hanania/output/movielens/")

#rename a column
final = final.withColumnRenamed("avg(rating)", "score")
```

```
#transform the genres column into a list
final = final.withColumn('genres',split(final.genres, '\\|'))

#create the sql view
final.createOrReplaceTempView("movielens")
```

Best movie per year

```
spark.sql(

“SELECT title, year, score FROM ( SELECT title, year, score,
ROW_NUMBER() OVER (PARTITION BY year ORDER BY score DESC) AS rank
FROM movielens) ranked_movies WHERE rank = 1”

).show(truncate=False)
```

This request contains multiple sub-requests. Firstly, for all year, we display the movies by their score descending, and using **row_number()**, we attribute an incremental value for all movies. So the movies that have the best score (mean of the marks given by the users) for each year, will have 1 for its ranking column. We could also use the number of votes to calculate the rank, but the mean of marks use it already.

After creating this column rank, we can select all the movies that have rank=1 and get this result :

title	year	score
.hack Liminality In the Case of Yuki Aihara	null	5.0
Gold Rush, The	1925	4.052878965922444
In the Meantime, Darling	1944	5.0
Geordie	1955	5.0
Robin Williams - Off the Wall	1978	5.0

Globally, the best movies always have 5 as a mean score. But for some years, like 1925, the best movies had just 4.05 as a mean score.

Furthermore, we can see that some movies don't have years, so their year column is **null**

Best movie par genre

```
spark.sql(

"SELECT * FROM ( SELECT title, genre, score, ROW_NUMBER() OVER
(PARTITION BY genre ORDER BY score DESC) as rank FROM movielens
LATERAL VIEW explode(genres) exploded_genre AS genre)
ranked_movies WHERE rank = 1"

).show(truncate=False)
```

This request is similar to the previous one. But, we also use the **explode** function on the **genre** column. Associated with the **lateral view**, This function will create multiple lines for every genre of a movie : title :[genreA, genreB] will create two lines title :A and title:B.

So now we can get the ranking with **row number()**, and select partition the data by genre.

title	genre	score	rank
Comin' Round the Mountain	Musical	5.0	1
A Sister's Revenge	Mystery	5.0	1
I Spy Returns	Action	5.0	1
The Laws of Thermodynamics	Romance	5.0	1

We can see that all movies have at least one genre.

For 'action' movie per year

```
spark.sql(

"SELECT * FROM (SELECT *, ROW_NUMBER() OVER (PARTITION BY year ORDER BY
score DESC) AS rank FROM ( SELECT title, year, score FROM movielens
LATERAL VIEW explode(genres) AS genre WHERE genre = 'Action' )
action_movies ) ranked WHERE rank = 1"

).show(truncate=False)
```

This request is a bit different from the previous one.

Here we have to preselect all the movies that have **Action** in their genre list, and then we can apply the same treatment as before, **row_number()** to select the best score partitioned by year

title	year	score	rank
Lazybones	1925	3.0	1
Between Two Worlds	1944	3.5714285714285716	1
Dam Busters, The	1955	3.727272727272727	1
Revenge	1978	5.0	1
Bedwin Hacker	2003	5.0	1
Ready Player One	null	3.6542986425339365	1
Fire!	1901	2.5	1

Best romance per year

```
park.sql("
SELECT * FROM (SELECT *, ROW_NUMBER() OVER (PARTITION BY year ORDER BY
score DESC) AS rank FROM ( SELECT title, year, score FROM movielens
LATERAL VIEW explode(genres) AS genre WHERE genre = 'Romance' ) romance)
ranked WHERE rank = 1
")
```

This request is exactly the same as the previous one. But we replace **action** by **romance**

title	year	score	rank
Parwaaz Hai Junoon	null	4.5	1
Gold Rush, The	1925	4.052878965922444	1
To Have and Have Not	1944	4.012206572769953	1
Geordie	1955	5.0	1
Sultan	1978	4.5	1
One Love	2003	5.0	1

We can see that the first movie doesn't have any year.

II. Bike rental model

This second exercise is about Bike rental predictions. We have a dataset that displays various information such as the weather, the hour, etc, and the number of demand per hour. Our objective is to predict the demand.

season	yr	mnth	hr	holiday	workingday	weathersit	temp	hum	windspeed	dayOfWeek	days	demand
1	0	1	0	0	0	1	0.24	0.81	0.0	Sat	0	16
1	0	1	1	0	0	1	0.22	0.8	0.0	Sat	0	40
1	0	1	2	0	0	1	0.22	0.8	0.0	Sat	0	32
1	0	1	3	0	0	1	0.24	0.75	0.0	Sat	0	13
1	0	1	4	0	0	1	0.24	0.75	0.0	Sat	0	1
1	0	1	5	0	0	2	0.24	0.75	0.0896	Sat	0	1
1	0	1	6	0	0	1	0.22	0.8	0.0	Sat	0	2

A. Model done in class

We have started this TP in class, so we start by analyzing its performances

Firstly, we transform the day into a category

```
StringIndexer(inputCol='dayOfWeek',  
outputCol='day_cat').fit(rowData).transform(rowData)
```

Secondly, we generate a vector of all the column

```
VectorAssembler(inputCols=  
['season', 'yr', 'mnth', 'hr', 'holiday', 'workingday', 'weathersit', 'temp', 'hum', 'windspeed', 'day_cat'], outputCol = 'features')
```

To predict the demand, we use the features column, and we split the dataframe into a train and a test sub dataset

```
modelData = data.select('features', 'demand')  
trainData, testData = modelData.randomSplit([0.7, 0.3])
```

features	demand	prediction
(11, [0, 1, 2, 6, 7, 8], [1.0, 1.0, 12.0, 2.0, 0.24, 0.7])	26.0	10.183484111148545
(11, [0, 1, 2, 6, 7, 8], [4.0, 1.0, 12.0, 1.0, 0.3, 0.7])	94.0	93.32759943722766
(11, [0, 2, 3, 6, 7, 8], [1.0, 1.0, 1.0, 1.0, 0.22, 0.8])	40.0	-78.30637991454373

Then we create a linear regression

```
lr = LinearRegression(labelCol='demand')
lrModel = lr.fit(trainData)
```

And then we evaluate the model on the train dataset

```
#test the model on the test Data
testResults = lrModel.evaluate(testData)
```

And the r2 score of the model on the test dataset is :

```
r2=0.382574
rootMeanSquaredError=142.383
```

So , the way we trained the model in this first part is very far from being satisfactory

We can get some insight of the prediction

Standard deviation and Average of the demand and the prediction by hour and seasons

```
#Hour and demand.
import pyspark.sql.functions as F
pred.groupBy("hr", "season").agg(

    F.avg("demand").alias("avg_real_demand"),
    F.avg("prediction").alias("avg_predicted_demand"),
    F.stddev("demand").alias("std_real_demand"),
    F.stddev("prediction").alias("std_pred_demand")

).sort('avg_real_demand').sort('hr').show()
```

hr	season	avg_real_demand	avg_predicted_demand	std_real_demand	std_predicted_demand
0	4	56.333333333333336	77.6407038651634	43.73352503653475	58.92857158923539
0	1	27.40449438202247	9.253572082479	21.243393975445947	67.69408214400802
0	2	56.84239130434783	82.3477241180681	40.34688586027427	68.39055150864472
0	3	73.9144385026738	138.93517519366523	45.24809928973331	56.380816806566855
1	2	35.22282608695652	83.42573246887578	31.9912781941294	67.11245774418246
1	1	18.140449438202246	12.444423695691297	19.078708333146423	66.3543807771044
1	3	43.72043010752688	142.2432432594821	37.77726059671858	54.56900030427858
1	4	35.92045454545455	80.93599306731419	36.6092015258956	57.76556260596105

We can see that depending of the hour, it is not always the same season that gets the highest demand

Standard deviation and Average of the demand and the prediction by working day

```
pred.groupBy("workingday").agg(  
    F.avg("demand").alias("avg_real_demand"),  
    F.avg("prediction").alias("avg_predicted_demand"),  
    F.stddev("demand").alias("std_real_demand"),  
    F.stddev("prediction").alias("std_predicted_demand")  
).show()
```

workingday	avg_real_demand	avg_predicted_demand	std_real_demand	std_predicted_demand
1	193.20775389801938	192.16539348556563	185.10747659123123	113.77161437615888
0	181.40533188248097	182.37639921666457	172.85383171939426	112.94808335764274

We can see that working day does not have an important impact on the demand


```
import pyspark.sql.functions as F  
pred.groupBy("weathersit").agg(  
    F.avg("demand").alias("avg_real_demand"),  
    F.avg("prediction").alias("avg_predicted_demand"),  
    F.stddev("demand").alias("std_real_demand"),  
    F.stddev("prediction").alias("std_predicted_demand")  
).show()
```

weathersit	avg_real_demand	avg_predicted_demand	std_real_demand	std_predicted_demand
1	204.8692718829405	207.55772726039677	189.48777342363275	114.18688398640381
3	111.57928118393235	129.3936522077394	133.78104501327016	101.7957886542607
4	74.33333333333333	24.03760423586199	77.92517778826901	75.3213319610337
2	175.16549295774647	161.33974746226298	165.43158927677524	103.03182924598701

But the weather has a important impact

Plotting the predictions and the real demand by temperature and weathersit

```
pred_ = pred.toPandas()
fig, ax = plt.subplots()
scatter = ax.scatter(pred_['temp'], pred_["demand"],
c=pred_["weathersit"])
ax.scatter(pred_['temp'], pred_["prediction"],
c=pred_["weathersit"], marker="x")
```



We can see that the temperature has an important impact on the demand. Indeed, there is 5 times more demand when it is 0.7 than is 0.

Furthermore, the weather situation is also important as the 1 and 2 are the most common.

B. Model improvement

To improve the model, we, we started by transform the days into a category, and we delete the **days** and **dayOfWeek** column as they are not useful.

```
# I transform all the data into numericals, and I drop the non
numericals
indexer = StringIndexer(inputCol='dayOfWeek', outputCol='day_cat')
indexed_data =indexer.fit(rowData).transform(rowData)

#removing the old dayOfWeek, and days as they are not usefull for
the
indexed_data = indexed_data.drop('dayOfWeek','days')
```

Then we use one hot encoding to add dummies variables

```
encoder = OneHotEncoder(  
inputCols=['season', 'yr', 'mnth', 'hr', 'holiday', 'workingday',  
'weathersit', 'day_cat'],  
outputCols=['seasonVec', 'yearVec', 'mnthVec', 'hrVec', 'holidayVec',  
'workingdayVec', 'weathersitVec', 'day_catVec']  
)  
  
#apply on the dataset  
model = encoder.fit(indexed_data)  
encoded_data = model.transform(indexed_data)  
#dropping the old column  
encoded_data = encoded_data.drop('season',  
'yr', 'mnth', 'hr', 'holiday', 'workingday', 'weathersit', 'day_cat')
```

So, all columns as like this

temp	hum	windspeed	demand	seasonVec	yearVec	mnthVec	hrVec	holidayVec
0.24	0.81	0.0	16	(4, [1], [1.0])	(1, [0], [1.0])	(12, [1], [1.0])	(23, [0], [1.0])	(1, [0], [1.0])
0.22	0.8	0.0	40	(4, [1], [1.0])	(1, [0], [1.0])	(12, [1], [1.0])	(23, [1], [1.0])	(1, [0], [1.0])

Then, we group all the columns together except the **demand**

```
col = ['seasonVec', 'yearVec', 'mnthVec', 'hrVec', 'holidayVec',  
'workingdayVec', 'weathersitVec', 'day_catVec', 'temp', 'hum', 'windspe  
ed']  
vector= VectorAssembler(inputCols=col, outputCol = 'features')
```

And we normalize the data,

```
#standard scaler  
standarScaler = StandardScaler(inputCol="features",  
outputCol="STfeatures")  
STdata = standarScaler.fit(data).transform(data)  
  
#minmaxscaler  
mmScaler = MinMaxScaler(inputCol="STfeatures",  
outputCol="NRfeatures")  
NRdata = mmScaler.fit(STdata).transform(STdata)
```

Now we have the data in a good format, we can create our model.

To create the model, we start by selecting the column we want to predict : **demand** and the column we want to use **NRfeatures** that represents the features normalized. We also split the dataset into a train and test sub datasets

```
modelData = NRdata.select('NRfeatures', 'demand')
trainData, testData = modelData.randomSplit([0.8, 0.2])
```

```
#create the linear regression
lr = LinearRegression(featuresCol="NRfeatures",labelCol='demand')

lrModel = lr.fit(trainData)
predictions = lrModel.transform(testData)

# we can see that the model is good !
print(lrModel.summary.r2)
lrModel.summary.meanAbsoluteError

r2 : 0.6885523286532669
MAE : 75.00301289062446
```

The results are pretty good on the training dataset!

We also use an evaluator to evaluate the model on the test dataset

```
lr_evaluator = RegressionEvaluator(predictionCol="prediction",
labelCol="demand",metricName="r2")

r2 : 0.6753477359544817
```

C. The second improvement

For the second improvement, we will use cross validation to test different models. Firstly we create a pipeline that will apply the linear regression

```
final_pipeline = Pipeline(stages=[lr])
```

And we also create a cross validator that will test different values for the model and find the best model

```
param_grid = ParamGridBuilder() \
    .addGrid(lr.regParam, [0.1, 0.01]) \
    .addGrid(lr.elasticNetParam, [0.0, 0.5, 1.0]) \
    .build()

cv = CrossValidator(estimator=final_pipeline,
    estimatorParamMaps=param_grid, evaluator=lr_evaluator, numFolds=5)
```

We fit it on the train dataset, and test it on the test dataset

```
# fit the cross-validator to the data
cv_model = cv.fit(trainData)

#transform on test data
tuned_lr_accuracy = lr_evaluator.evaluate(cv_model.
    transform(testData))

tuned_lr_mae = lr_evaluator.evaluate(cv_model.transform(testData),
    {lr_evaluator.metricName: "mae"})

tuned logistic regression accuracy :0.676911332147188
tuned logistic regression mae:74.5997047411219
```


D. Comparison with other machine learning model

After obtaining this first result, we decide to compare its score with other machine learning models

```
# create the models Random forest and decision tree
dt = DecisionTreeRegressor(featuresCol="NRfeatures",
labelCol="demand")

rf = RandomForestRegressor(featuresCol="NRfeatures",
labelCol="demand")

#The pipelines
dt_pipeline = Pipeline(stages=[dt])

rf_pipeline = Pipeline(stages=[rf])

# create the parameter grids for each model
dt_param_grid = ParamGridBuilder() \
    .addGrid(dt.maxDepth, [5, 10,15,20]) \
    .build()

rf_param_grid = ParamGridBuilder() \
    .addGrid(rf.maxDepth, [5, 10,15,20]) \
    .addGrid(rf.numTrees, [10, 20]) \
    .build()

# set up the cross-validators for each model
dt_cv = CrossValidator(estimator=dt_pipeline,
estimatorParamMaps=dt_param_grid, evaluator=lr_evaluator,
numFolds=5)

rf_cv = CrossValidator(estimator=rf_pipeline,
estimatorParamMaps=rf_param_grid, evaluator=lr_evaluator,
numFolds=5)

# fit the cross-validators to the data and evaluate the best
models
dt_cv_model = dt_cv.fit(trainData)
rf_cv_model = rf_cv.fit(trainData)
```

```
#get best models
best_dt_model = dt_cv_model.bestModel
best_rf_model = rf_cv_model.bestModel

#predictions
dt_predictions = best_dt_model.transform(testData)
rf_predictions = best_rf_model.transform(testData)

#evaluation
dt_mae lr_evaluator.evaluate(dt_predictions,
{lr_evaluator.metricName: "mae"})

dt_r2 = lr_evaluator.evaluate(dt_predictions)
```

Finally, we can compare the different models. It is the random forest model that has the best accuracy score.

III. SMS Spam Collection

A. Data Preparation

For this exercise, we have a complete dataset. It represents some SMS and explains if they are spam or safe (ham). The objective is to predict if a SMS is safe or not.

```
+-----+-----+
|label|text
+-----+-----+
|ham  |Go until jurong point, crazy.. Available only in bugis
|ham  |Ok lar... Joking wif u oni...
+-----+-----+
```

The dataset is not in a good format to make predictions on it, we need to make transformations.

The first is to transform the label into categories. We use **StringIndexer** to transform "ham" on 0, and "spam" on 1

```
indexer = StringIndexer(inputCol="label", outputCol="labelIndex")
```

We also transform the text into vectors

```
tokenizer = Tokenizer(inputCol="text", outputCol="sms_words")
```

```
+-----+-----+-----+-----+
|label|          text|labelIndex|          sms_words|
+-----+-----+-----+-----+
| ham|Go until jurong p...|      0.0|[go, until, juron...|
| ham|Ok lar... Joking ...|      0.0|[ok, lar..., joki...|
+-----+-----+-----+-----+
```

Then, we remove stop words because they are not useful for the prediction

```
remover = StopWordsRemover(inputCol="sms_words",
outputCol="filtered_sms_words")
```

And we transform this list of words into list of number

CountVectorizer is used to extract a vocabulary from document collections.

```
count_vectorizer = CountVectorizer(inputCol="filtered_sms_words",
outputCol="raw_features")
```

label	text labelIndex	sms_words	filtered_sms_words	raw_features
ham	Go until jurong p...	0.0 [go, until, juron...	[go, jurong, poin...	(13423,[7,11,31,6...
ham	Ok lar... Joking ...	0.0 [ok, lar..., joki...	[ok, lar..., joki...	(13423,[0,24,301,...

In this dataset, the **raw_features** represents the text. But the different words don't have an importance indicator. That is why we will use the TDIDF transformer, to get the most important words in spam messages.

```
idf = IDF(inputCol="raw_features", outputCol="features")
```

Finally, we merge multiple columns into a vector column using **VectorAssembler**

```
assembler = VectorAssembler(inputCols=["features"],
outputCol="features_vector")
```

All the predictions will be done with this **features_vectors**

B. Training 4 models

After the data preparation, we obtain a vector with numerical values. This vector is in a good format to make predictions.

We will use 4 different models : a **Logistic Regression**, a Decision Tree, a **Random Forest**, and a **Naive Bayes** to see which one is the best

```
#Creating pipelines for the different classifiers

lr = LogisticRegression(featuresCol="features_vector",
labelCol="labelIndex")
pipeline_lr = Pipeline(stages=[assembler, lr])
```

```

dt = DecisionTreeClassifier(featuresCol="features_vector",
labelCol="labelIndex")
pipeline_dt = Pipeline(stages=[assembler, dt])

rf = RandomForestClassifier(featuresCol="features_vector",
labelCol="labelIndex")
pipeline_rf = Pipeline(stages=[assembler, rf])

nb = NaiveBayes(featuresCol="features_vector",
labelCol="labelIndex")
pipeline_nb = Pipeline(stages=[assembler, nb])

```

For all of these models, we use a pipeline, they are used to combine several data processing steps (the assembler and the model creation) into one well organized sequence.

Then, we need to train our models. So we separate the whole dataset into a train and a test part.

```

#Separate the data into a train and a test data
training, testing = df.randomSplit([0.8, 0.2])

#Training the models using the pipelines
lr_model = pipeline_lr.fit(training)
dt_model = pipeline_dt.fit(training)
rf_model = pipeline_rf.fit(training)
nb_model = pipeline_nb.fit(training)

```

Then we define an evaluator that will give a score for each model

```

#Evaluating the models on the testing data
evaluator =
MulticlassClassificationEvaluator(predictionCol="prediction",
labelCol="labelIndex", metricName="accuracy")

#Calculating the accuracies
lr_accuracy = evaluator.evaluate(lr_model.transform(testing))
dt_accuracy = evaluator.evaluate(dt_model.transform(testing))
rf_accuracy = evaluator.evaluate(rf_model.transform(testing))
nb_accuracy = evaluator.evaluate(nb_model.transform(testing))

```

```
Logistic Regression Accuracy: 0.9836065573770492
Decision Tree Accuracy: 0.9326047358834244
Random Forest Accuracy: 0.8752276867030966
Naive Bayes Accuracy: 0.9153005464480874
```

We can see that the best model is the Logistic Regression and the least is the random forest.

So, for the next part, we will try to improve the random forest model.

C. Tuning the Random Forest model

In order to improve the random forest model, we will change some hyperparameters. We try some tree numbers values (10,20,50) and some maxDepth(5,10,15).

And we use a **CrossValidator** to test the different hyperparameters.

```
param_grid = (ParamGridBuilder().addGrid(rf.numTrees, [10, 20, 50]).addGrid(rf.maxDepth, [5, 10, 15]).build())

#CrossValidator
crossval = CrossValidator(estimator=pipeline_rf,
estimatorParamMaps=param_grid, evaluator=evaluator, numFolds=4)
rf_model_tuned = crossval.fit(training)
```

The new random forest accuracy is

```
Tuned Random Forest Accuracy: 0.9180327868852459
```

The final results are :

Logistic Regression Accuracy	0.9836065573770492
Decision Tree Accuracy	0.9326047358834244
Tuned Random Forest Accuracy	0.9180327868852459
Naive Bayes Accuracy	0.9153005464480874
Random Forest Accuracy	0.8752276867030966

By tuning the random forest, we get a 5% improvement.